




Background of Cloud computing

1. Distributed Systems



❖ What is a distributed system?

Distributed System: a group of independent/autonomous computers that are networked together appear to the user as a one computer.

- Work together to achieve a common goal.
- Clouds can be DS ?



-- History --

- Problems that are larger than what a single machine can handle
- Computer Networks, Message passing were invented to facilitate distributed systems
- ARPANET eventually became Internet

What is a distributed system?

A collection of independent computers

A communication facility to pass messages

No shared memory

No shared clock

Each computer has its own operating system



Where they are used?

- **Strategic Systems (Defense/Intelligence)**
- **Bioinformatics (Genome/ Molecular Analysis)**
- **Visualization and Graphics (VR)**
- **Economics and Finance (Fin-Tech)**
- **Scientific Computing (Weather forecasting)**

Advantages and Disadvantages

Advantages

- Communication and resource sharing possible
- Economics – price-performance ratio
- Reliability, scalability
- Potential for incremental growth

Disadvantages

- Distribution-aware PLs, OSs and applications
- Network connectivity essential
- Security and privacy



Design Goals/ Characteristics

- **Distributed Transparency**

- Location

- Migration

- Replication

- Concurrency

- **Openness**

- **Scalability**

- **Fault-tolerance**

- **High availability**

- **Recoverability**

- **Performance Predictability**

- **Security**

Transparency in a Distributed System

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource may be shared by several competitive users
Concurrency	Hide that a resource may be performed by several heterogeneous resources(h/w)
Failure	Hide the failure and recovery of a resource
Persistence	Hide whether a (software) resource is in memory or on disk

Parallel vs. Distributed Systems

A concurrent system could be Parallel or Distributed:

Two possible Views to make the distinction

.

View 1:

.

- Parallel System: a particular tightly-coupled form of distributed computing
- Distributed System: a loosely-coupled form of parallel computing

.

View 2:

.

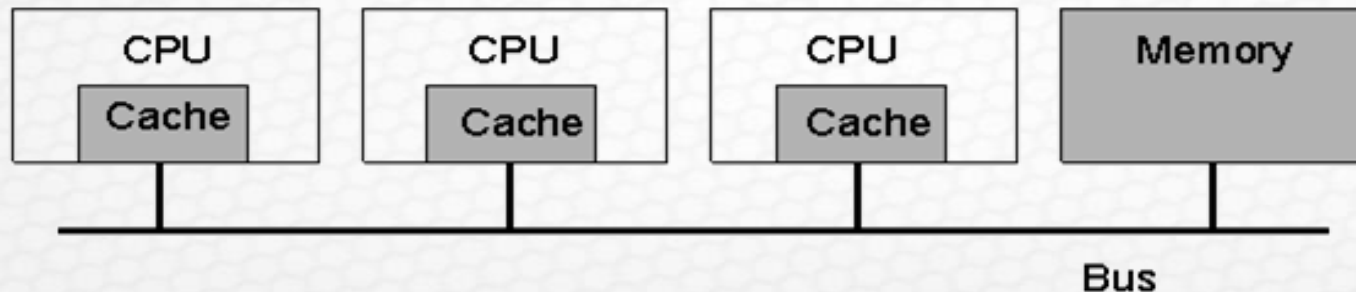
- Parallel System: processors access a shared memory to exchange information
- Distributed System: uses a "distributed memory". Message passing is used to exchange information between the processors as each one has its own private memory.

❖ Hardware Concepts: Multiprocessors(1/2)

■ Multiprocessor dimensions

- Memory: could be shared or be private to each CPU
- Interconnect: could be shared (bus-based) or switched


■ A bus-based multiprocessor



Hardware Concepts: Multiprocessors(2/2)

Memories

CPUs




Crosspoint switch

(a) A crossbar switch

CPUs


Memories



2x2 switch


(b) An omega switching network

Homogeneous Multicomputer Systems



(a)

(a) Grid




(b)

(b) Hypercube

Hardware Model

- **The four possible combinations: (Flynn's Taxonomy)**
 - SISD**: in traditional uniprocessor computers
 - MISD**: Multiple concurrent instructions operating on the same data element. Not useful!
 - SIMD**: Single instruction operates on multiple data elements in parallel
 - MIMD**: Covers parallel & distributed systems and machines that contain multiple computers
- **Distributed systems are MIMD**





Distributed Systems Models

■ Minicomputer model (e.g., early networks)

- Each user has local machine
- Local processing but can fetch remote data (files, databases)

■ Workstation model (e.g., Sprite)

- Processing can also migrate

■ Client-server Model (e.g., V system, world wide web)

- User has local workstation
- Powerful workstations serve as servers (file, print, DB servers)

■ Processor pool model (e.g., Amoeba, Plan 9)

- Terminals are Xterms or diskless terminals
- Pool of backend processors handle processing

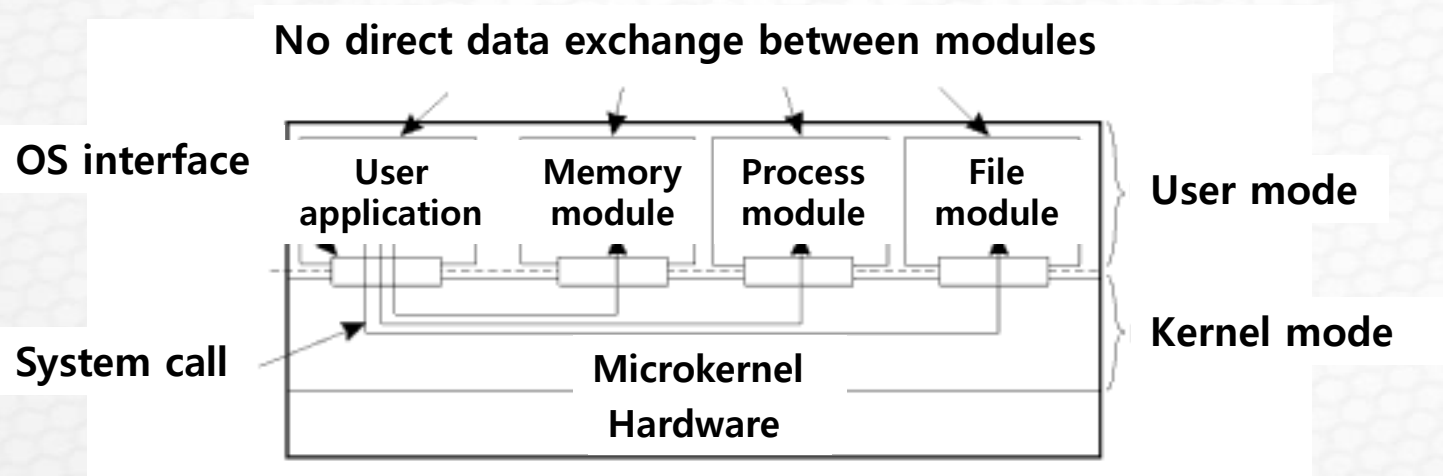
Uniprocessor Operating Systems(1/2)

- **An OS acts as a resource manager or an arbitrator**
 - **Manages CPU, I/O devices, memory**
- **OS provides a virtual interface that is easier to use than hardware**
- **Structure of uniprocessor operating systems**
 - **Monolithic (e.g., MS-DOS, early UNIX)**
 - ✓ **One large kernel that handles everything**
 - **Layered design**
 - ✓ **Functionality is decomposed into N layers**
 - ✓ **Each layer uses services of layer N-1 and implements new service(s) for layer N+1**

Uniprocessor Operating Systems(2/2)

Microkernel architecture

- Small kernel
- user-level servers implement additional functionality





Distributed Operating System

- Manages resources in a distributed system
 - Seamlessly and transparently to the user
- Looks to the user like a centralized OS
 - But operates on multiple independent CPUs
- Provides transparency
 - Location, migration, concurrency, replication,...
- Presents users with a virtual uniprocessor

Types of Distributed OSs

System	Description	Main Goal
DOS	Tightly-coupled operating system for multi-processors and homogeneous multicomputers	Hide and manage hardware resources
NOS	Loosely-coupled operating system for heterogeneous multicomputers (LAN and WAN)	Offer local services to remote clients
Middleware	Additional layer atop of NOS implementing general-purpose services	Provide distribution transparency

Multiprocessor Operating Systems (1/2)

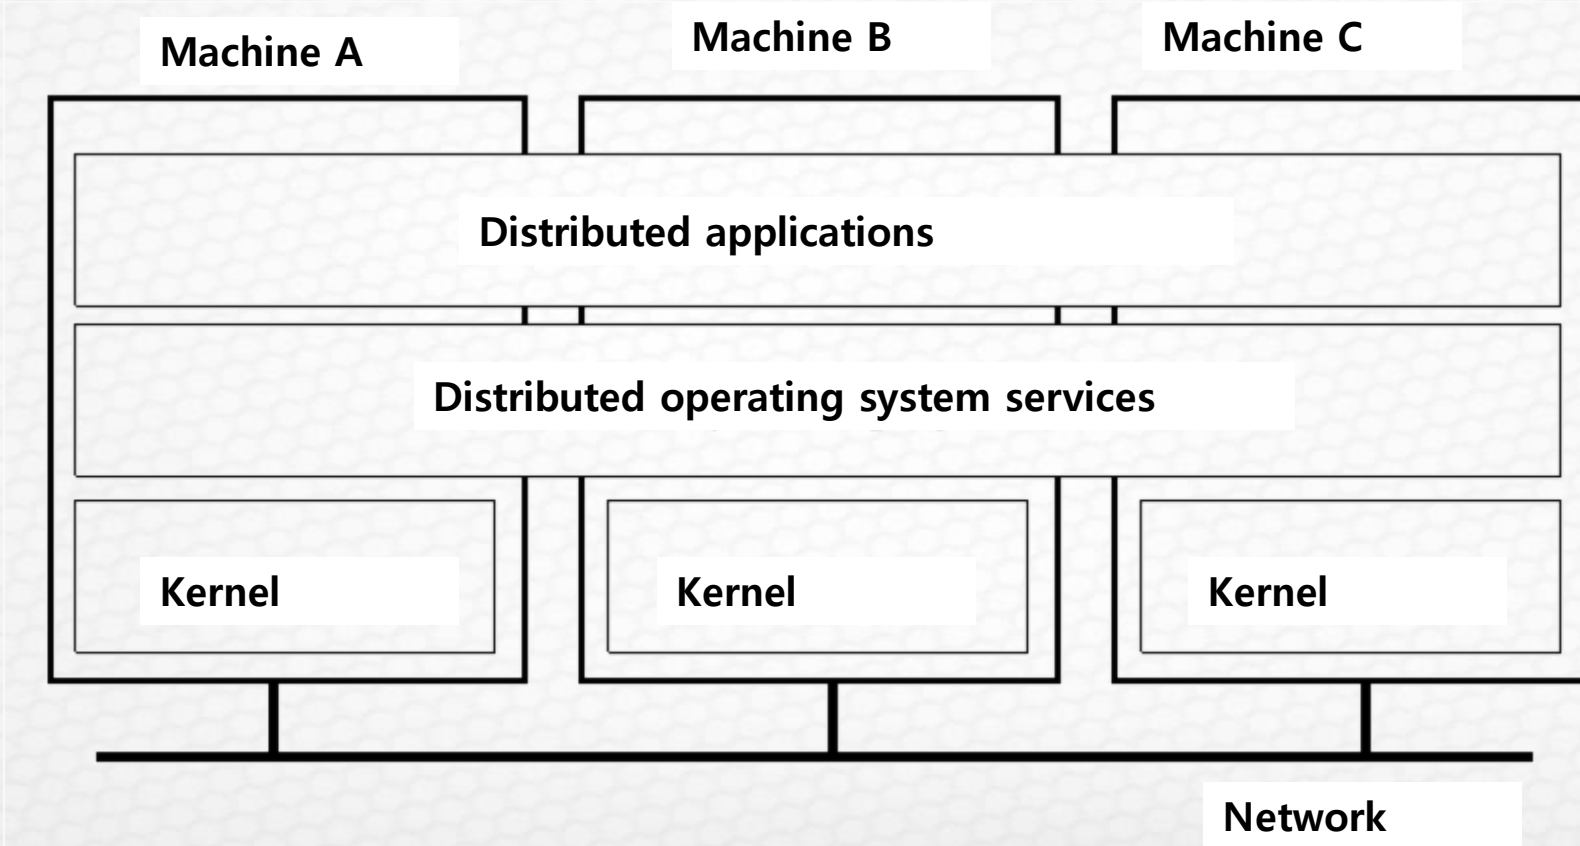
Like a uniprocessor operating system

Manages multiple CPUs transparently to the user

Each processor has its own hardware cache

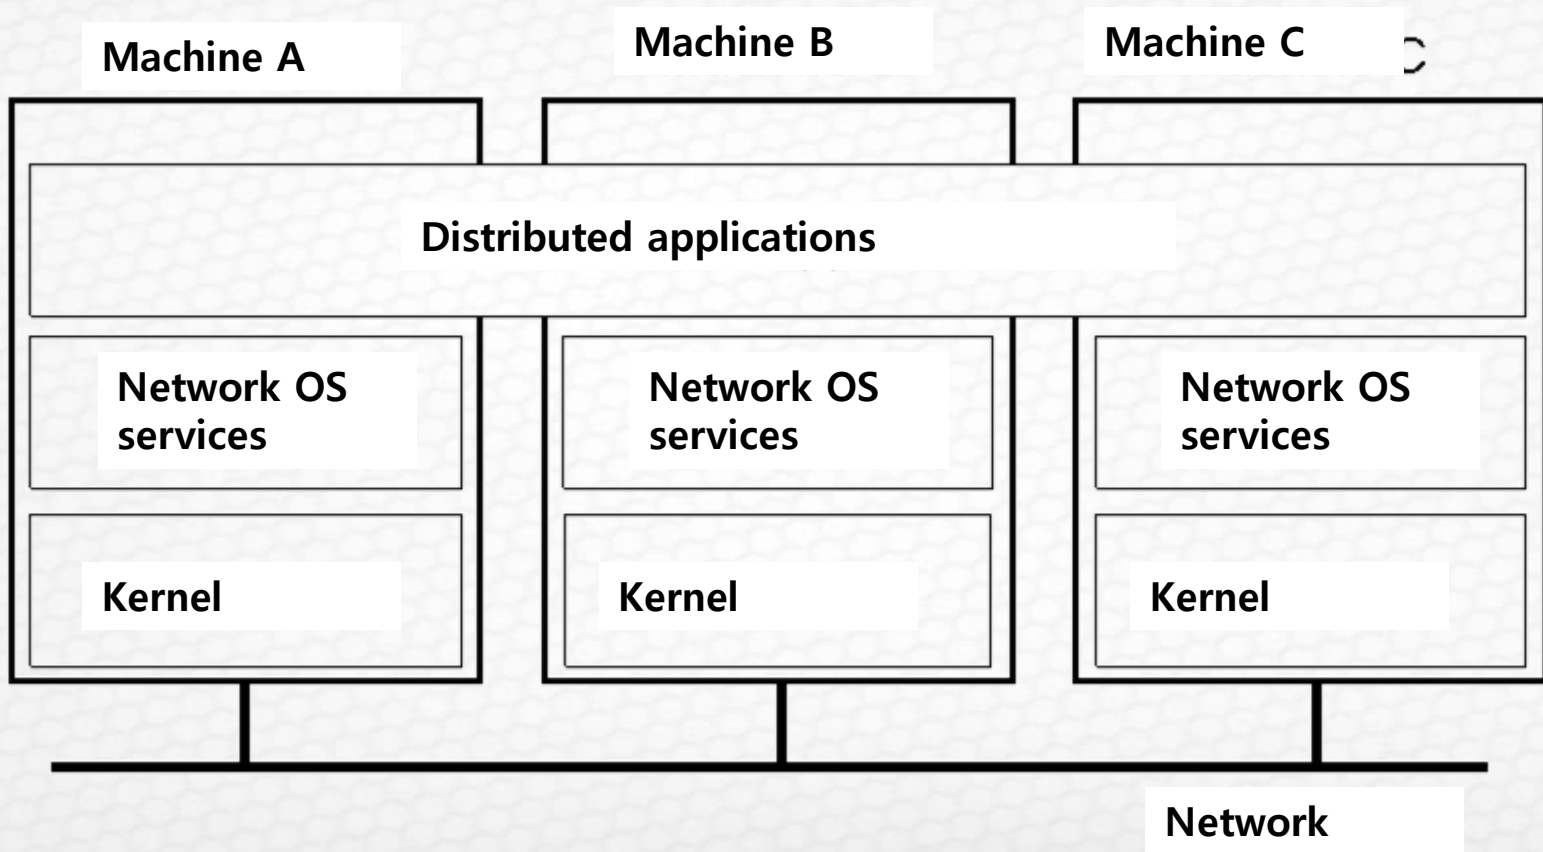
➡ Maintain consistency of cached data

Multicomputer Operating Systems (2/2)






Network Operating System (1/2)



Network Operating System (2/2)

■ Employs a client-server model

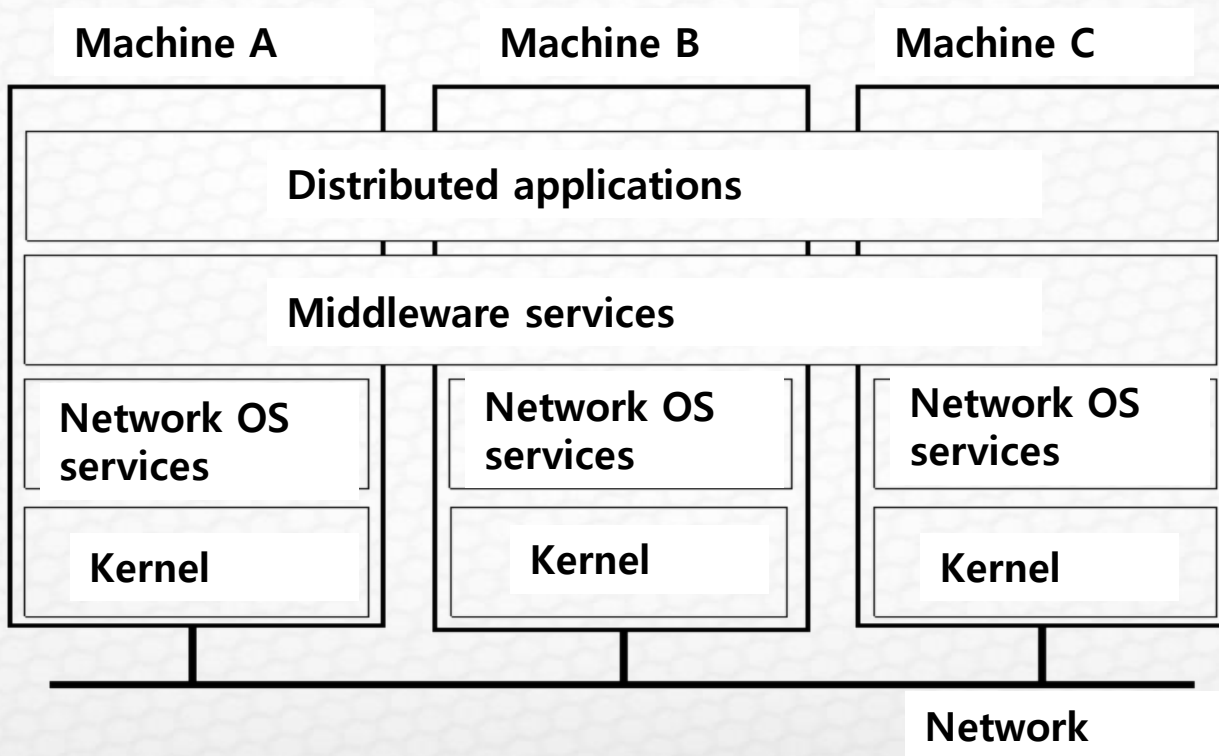
- Minimal OS kernel
- Additional functionality as user processes





Middleware-based Systems

■ General structure of a distributed system as middleware



Distributed system organizations

Microcomputer model

➡ several multiuser systems

Workstations/PCs model

➡ each user has own WS/PC to do work

➡ each user shares files and other resources

Processor pool model

LANs, MANs, WANs, WWW




Distributed System : HW vs SW


Hardware		Software	
		Loosely-coupled	Tightly-coupled
	Loosely-coupled	<ul style="list-style-type: none">■ Software Type: Network OS■ Multicomputer■ Each machine running its OS. Oses may differ.	<ul style="list-style-type: none">■ Software Type: Integrated Distributed System■ Group of Shared machines work like one computer but do not have shared memory
	Tightly-coupled	<ul style="list-style-type: none">■ Doesn't make sense	<ul style="list-style-type: none">■ Software Type: Multiprocessor Timesharing System■ E.g. UNIX machine with several processors and several terminals



Connection Topology



(a) Physical view




(b) Logical view

Computers in a Networked Environment. (*Galli, p. 3*)




Connection Topology



Connecting LAN Subnets with a Backbone. (*Galli, p.6*)



Connection Topology




Common Wired LAN Topologies. (*Galli, p.7*)


Six Computing Paradigm

Phases


1. Mainframe computing




2. PC computing




3. Network computing




4. Internet computing



5. Grid computing



6. Cloud computing



It looks like that cloud computing is a return to the **original mainframe** computing paradigm

However, there are several important differences

- Mainframe computing offers **finite**
- **Computing power**
- Cloud computing provides **almost infinite power and capacity**
- In mainframe computing, **dummy terminals** acted as user interface devices
- In cloud computing, **powerful PCs** can provide local computing power and caching support

2. Issues in Distributed Systems





Unable to determine up-to-date global state

- ➡ **no global memory**
- ➡ **no common clock**
- ➡ **unpredictable message delays**

Need device-efficient distributed control

- ➡ **e.g. how to get a consensus**



Need method for ordering events

Naming

All objects are named

Need to map name onto its location

Need a directory (or directories)

-  **replicated (to maintain consistency)**
-  **partitioned (which partition helps me?)**

Scalability

- ➡ Can system grow without performance degradation?
- ➡ Want to avoid centralized components

Process synchronization

- ➡ Enforce mutual exclusion to shared resources
- ➡ Deal with potential for deadlock

Compatibility

- Possible at different levels
- Binary level: all processing elements run same binary code
- Execution level: same source code can be compiled and run on all nodes
- Protocol level: all processing elements support same protocols

Data migration: bring data to the location

- distributed file system
- distributed shared memory

Computation migration

- e.g. RPC
- e.g. send a query for info computed remotely instead of requesting raw data

Distributed scheduling

- process migration

Authentication

→ **verify user identification**

Authorization

→ **determine user privileges**